


# Kalman Filter

16-385 Computer Vision (Kris Kitani)  
**Carnegie Mellon University**

Examples up to now have been **discrete** (binary) random variables

Kalman 'filtering' can be seen as a special case of a temporal inference with continuous random variables



Everything is continuous...

$x$        $e$        $P(x_0)$        $P(e|x)$        $P(x_t|x_{t-1})$

probability distributions are no longer tables but functions

# Making the connection to the ‘filtering’ equations

## (Discrete) Filtering

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \sum_{\mathbf{X}_t} P(\mathbf{X}_{t+1} | \mathbf{X}_t) P(\mathbf{X}_t | \mathbf{e}_{1:t})$$

Tables    Tables    Tables

## Kalman Filtering

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \int_{\mathbf{x}_t} P(\mathbf{X}_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{e}_{1:t}) d\mathbf{x}_t$$


Gaussian    Gaussian    Gaussian


observation model    motion model    belief

integral because  
continuous PDFs

Simple, 1D example...

$x$






$$x_t = x_{t-1} + s + r_t$$

$$r_t \sim \mathcal{N}(0, \sigma_R)$$


'sampled from'

# System (motion) model



*How do you represent the motion model?*

$$P(x_t | x_{t-1})$$



*How do you represent the motion model?*


A linear Gaussian (continuous) transition model

$$P(x_t | x_{t-1}) = \mathcal{N}(x_t; x_{t-1} + s, \sigma_r)$$

mean

standard deviation

*How can you visualize this distribution?*

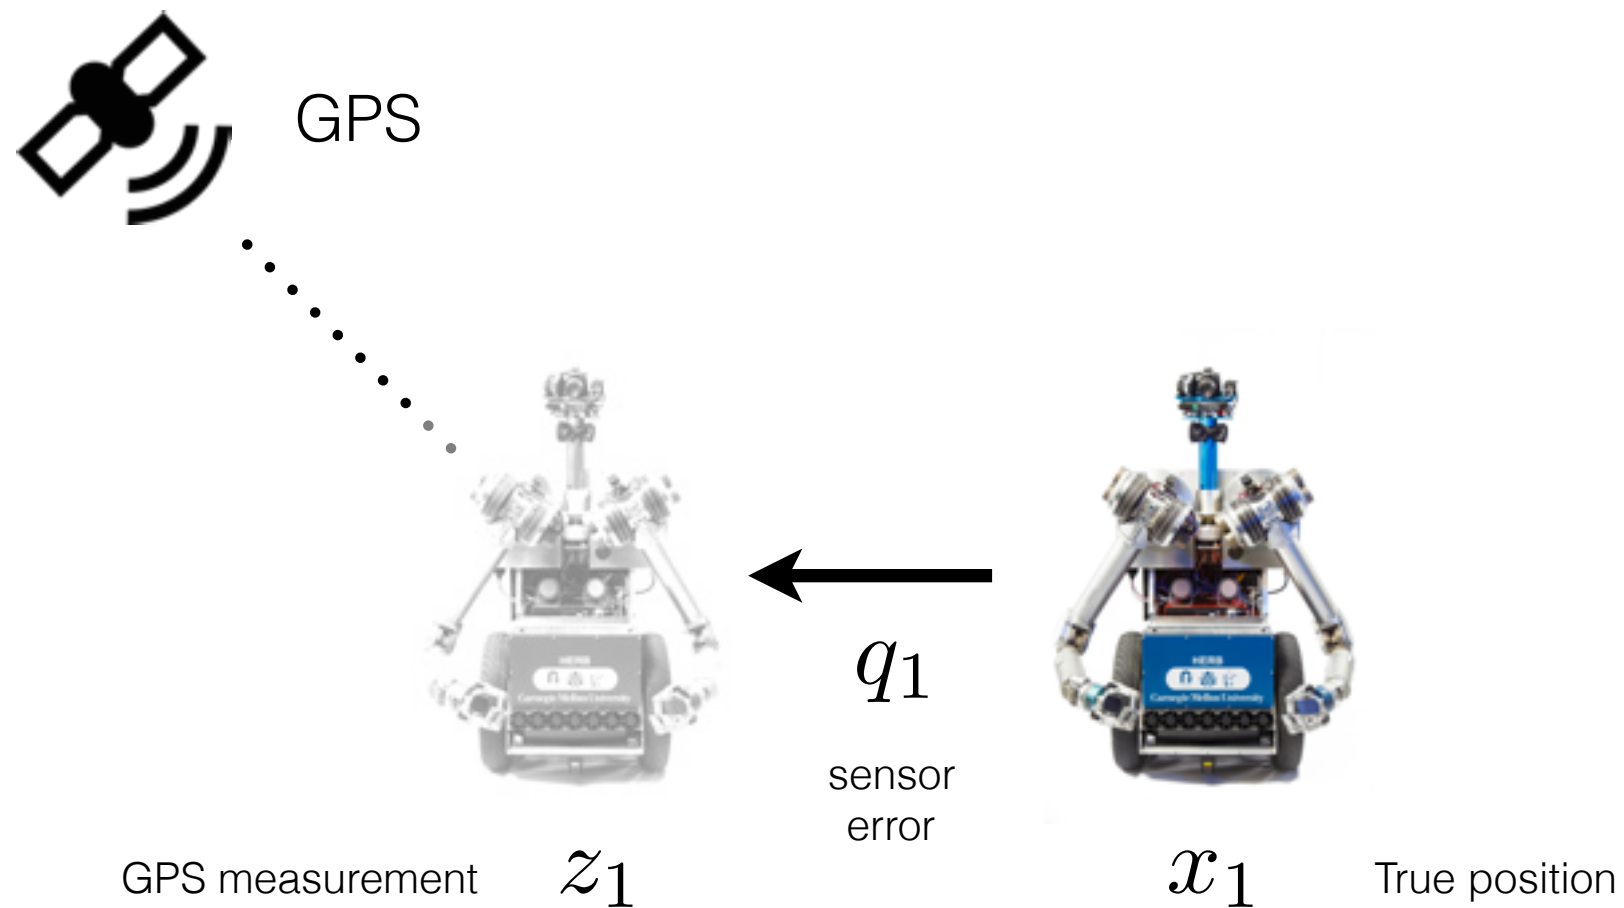


A linear Gaussian (continuous) transition model

$$P(x_t | x_{t-1}) = \mathcal{N}(x_t; x_{t-1} + s, \sigma_r)$$

*Why don't we just use a table as before?*






$$z_t = x_t + q_t$$

$$q_t \sim \mathcal{N}(0, \sigma_Q)$$

sampled from a Gaussian


# Observation (measurement) model



*How do you represent the observation (measurement) model?*

$$P(e|x)$$

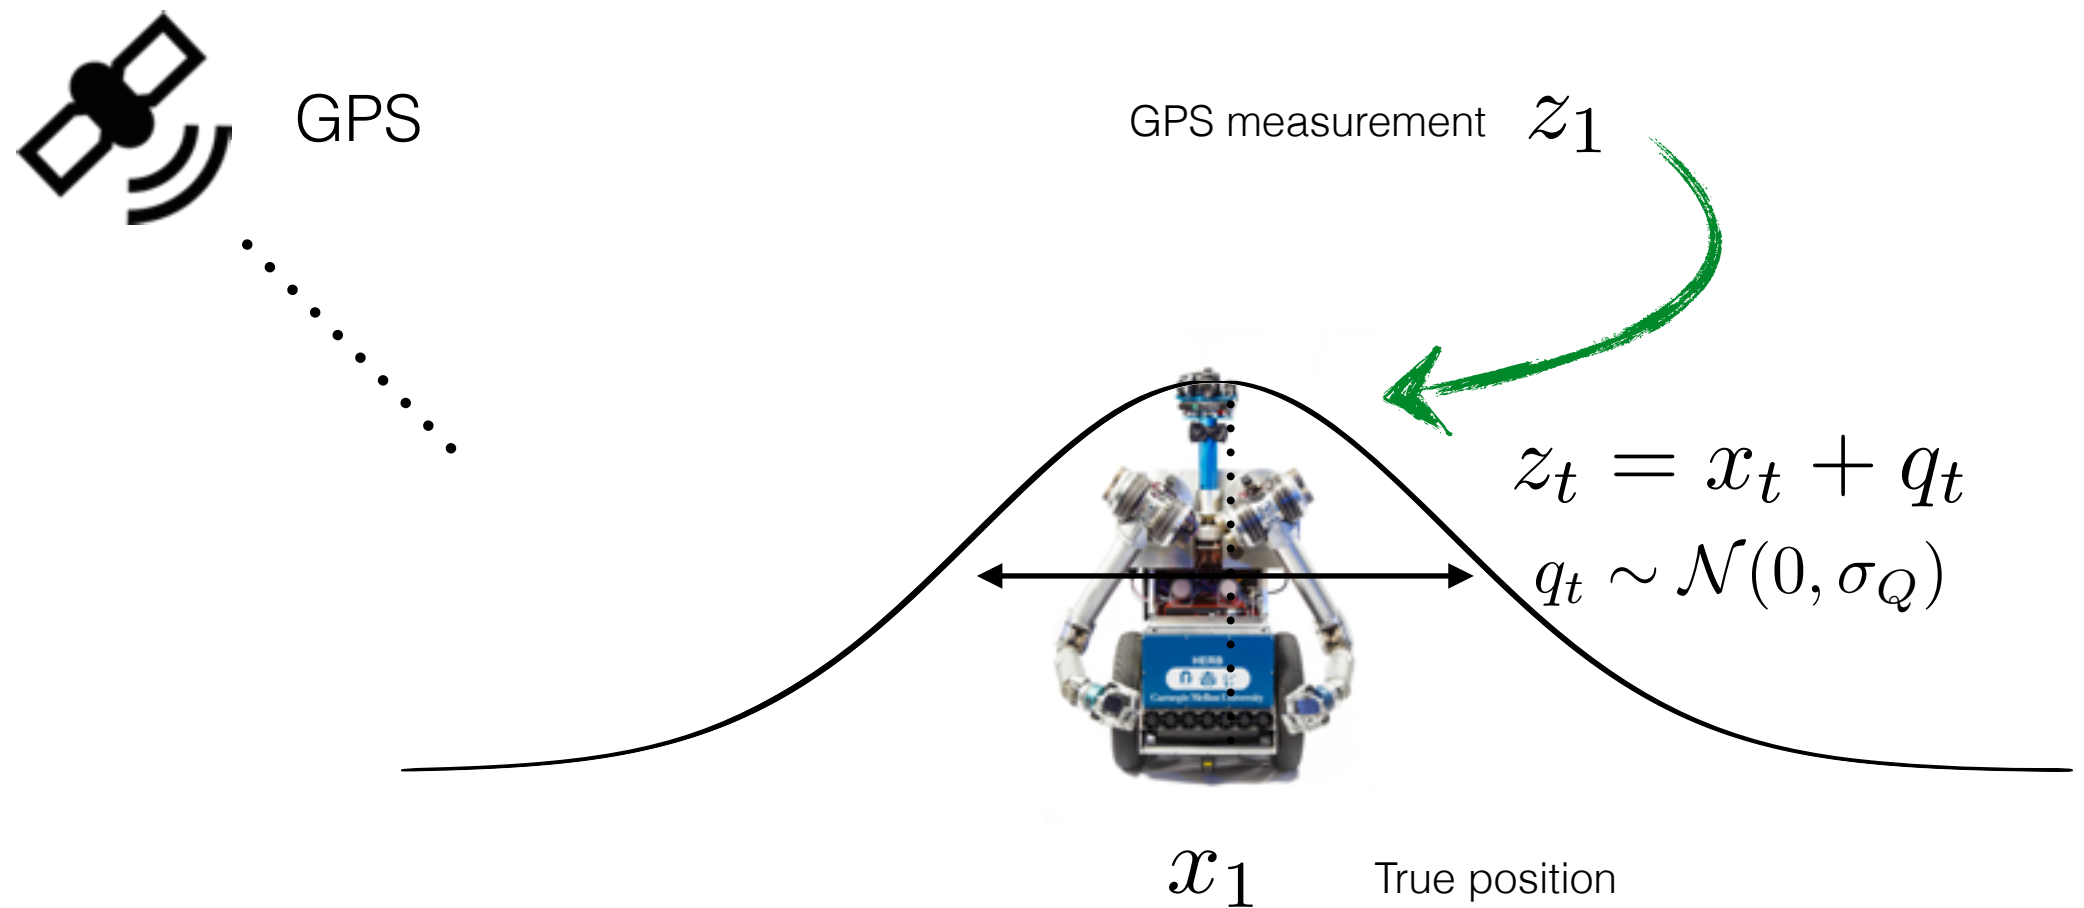
e represents z



*How do you represent the observation (measurement) model?*

Also a linear Gaussian model


$$P(z_t | x_t) = \mathcal{N}(z_t; x_t, \sigma_Q)$$



*How do you represent the observation (measurement) model?*

Also a linear Gaussian model

$$P(z_t | x_t) = \mathcal{N}(z_t; x_t, \sigma_Q)$$



$x_0$


true position

$\hat{x}_0$

initial estimate

initial estimate uncertainty  $\sigma_0$

# Prior (initial) State




$x_0$

true position

$\hat{x}_0$

initial estimate


*How do you represent the prior state probability?*



*How do you represent the prior state probability?*

Also a linear Gaussian model!

$$P(\hat{x}_0) = \mathcal{N}(\hat{x}_0; x_0, \sigma_0)$$



*How do you represent the prior state probability?*

Also a linear Gaussian model

$$P(\hat{x}_0) = \mathcal{N}(\hat{x}_0; x_0, \sigma_0)$$



# Inference

*So how do you do temporal filtering with the KL?*

**Recall:** the first step of filtering was the ‘prediction step’

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \int_{\mathbf{x}_t} \underbrace{P(\mathbf{X}_{t+1} | \mathbf{x}_t)}_{\text{motion model}} \underbrace{P(\mathbf{x}_t | \mathbf{e}_{1:t})}_{\text{belief}} d\mathbf{x}_t$$

prediction step

compute this!  
It's just another Gaussian



need to compute the ‘prediction’ mean and variance...

# Prediction

(Using the motion model)

*How would you predict  $\hat{x}_1$  given  $\hat{x}_0$  ?*

using this 'cap' notation to  
denote 'estimate'

$$\hat{x}_1 = \hat{x}_0 + s \quad (\text{This is the mean})$$

$$\sigma_1^2 = \sigma_0^2 + \sigma_r^2 \quad (\text{This is the variance})$$

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \int_{\mathbf{x}_t} \underbrace{P(\mathbf{X}_{t+1} | \mathbf{x}_t)}_{\text{motion model}} \underbrace{P(\mathbf{x}_t | \mathbf{e}_{1:t})}_{\text{belief}} d\mathbf{x}_t$$

prediction step

the second step after prediction is ...


... update step!

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto \underbrace{P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1})}_{\text{observation}} \int_{\mathbf{x}_t} \underbrace{P(\mathbf{X}_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{e}_{1:t})}_{\text{prediction}} d\mathbf{x}_t$$

compute this  
(using results of the prediction step)



In the **update step**, the **sensor measurement** **corrects** the system **prediction**



*Which estimate is correct? Is there a way to know?*

*Is there a way to merge this information?*

***Intuitively***, the smaller variance mean less uncertainty.



system  
prediction  $\sigma_1^2$



sensor  
estimate  $\sigma_q^2$

So we want a weighted state estimate correction



something  
like this...

$$\hat{x}_1^+ = \frac{\sigma_q^2}{\sigma_1^2 + \sigma_q^2} \hat{x}_1 + \frac{\sigma_1^2}{\sigma_1^2 + \sigma_q^2} z_1$$

***This happens naturally in the Bayesian filtering (with Gaussians) framework!***

Recall the filtering equation:

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto \overbrace{P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1})}^{\text{observation}} \overbrace{\int_{\mathbf{x}_t} P(\mathbf{X}_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{e}_{1:t}) d\mathbf{x}_t}^{\text{one step motion prediction}}$$

 Gaussian                       Gaussian

*What is the product of two Gaussians?*



Recall ...

When we multiply the prediction (Gaussian) with the observation model (Gaussian) we get ...

... a product of two Gaussians

$$\mu = \frac{\mu_1 \sigma_2^2 + \mu_2 \sigma_1^2}{\sigma_2^2 + \sigma_1^2}$$

$$\sigma = \frac{\sigma_1^2 \sigma_2^2}{\sigma_1^2 + \sigma_2^2}$$

applied to the filtering equation...

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \int_{\mathbf{x}_t} P(\mathbf{X}_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{e}_{1:t}) d\mathbf{x}_t$$

mean:  $z_1$   
variance:  $\sigma_q$

mean:  $\hat{x}_1$   
variance:  $\sigma_1$

new mean:

new variance:

$$\hat{x}_1^+ = \frac{\hat{x}_1 \sigma_q^2 + z_1 \sigma_1^2}{\sigma_q^2 + \sigma_1^2}$$

$$\hat{\sigma}_1^{2+} = \frac{\sigma_q^2 \sigma_1^2}{\sigma_q^2 + \sigma_1^2}$$

'plus' sign means post  
'update' estimate



system  
prediction  $\sigma_1^2$



sensor  
estimate  $\sigma_q^2$

With a little algebra...

$$\hat{x}_1^+ = \frac{\hat{x}_1 \sigma_q^2 + z_1 \sigma_1^2}{\sigma_q^2 + \sigma_1^2} = \hat{x}_1 \frac{\sigma_q^2}{\sigma_q^2 + \sigma_1^2} + z_1 \frac{\sigma_1^2}{\sigma_q^2 + \sigma_1^2}$$

We get a weighted state estimate correction!

# Kalman gain notation

With a little algebra...

$$\hat{x}_1^+ = \hat{x}_1 + \frac{\sigma_1^2}{\sigma_q^2 + \sigma_1^2} (z_1 - \hat{x}_1) = \hat{x}_1 + \underset{\substack{\uparrow \\ \text{'Kalman gain'}}}{K} (\underset{\substack{\uparrow \\ \text{'Innovation'}}}{z_1 - \hat{x}_1})$$

With a little algebra...

$$\sigma_1^+ = \frac{\sigma_1^2 \sigma_q^2}{\sigma_1^2 + \sigma_q^2} = \left( 1 - \frac{\sigma_1^2}{\sigma_1^2 + \sigma_q^2} \right) \sigma_1^2 = (1 - \mathbf{K}) \sigma_1^2$$

# Summary (1D Kalman Filtering)

To solve this...

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \int_{\mathbf{x}_t} P(\mathbf{X}_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{e}_{1:t}) d\mathbf{x}_t$$

Compute this...

$$\hat{x}_1^+ = \hat{x}_1 + \frac{\sigma_1^2}{\sigma_1^2 + \sigma_q^2} (z_1 - \hat{x}_1) \quad \sigma_1^{2+} = \sigma_1^2 - \frac{\sigma_1^2}{\sigma_1^2 + \sigma_q^2} \sigma_1^2$$

$$K = \frac{\sigma_1^2}{\sigma_1^2 + \sigma_q^2}$$

‘Kalman gain’

$$\hat{x}_1^+ = \hat{x}_1 + K(z_1 - \hat{x}_1)$$

mean of the new Gaussian

$$\sigma_1^{2+} = \sigma_1^2 - K \sigma_1^2$$

variance of the new Gaussian

# Simple 1D Implementation

$$[x \ p] = KF(x, v, z)$$

$$x = x + s;$$

$$v = v + q;$$

$$K = v / (v + r);$$

$$x = x + K * (z - x);$$

$$p = v - K * v;$$

Just 5 lines of code!

or just 2 lines

$$\begin{aligned} [x \ P] &= KF(x, v, z) \\ x &= (x+s) + (v+q) / ((v+q)+r) * (z - (x+s)) ; \\ p &= (v+q) - (v+q) / ((v+q)+r) * v ; \end{aligned}$$

# Bare computations (algorithm) of Bayesian filtering:

KalmanFilter( $\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$ )

prediction mean  $\bar{\mu}_t = A_t \overset{\text{motion}}{\mu_{t-1}} + B \overset{\text{control}}{u_t}$  'old' mean Prediction

prediction covariance  $\bar{\Sigma}_t = A_t \overset{\text{'old' covariance}}{\Sigma_{t-1}} A_t^\top + R$  Gaussian noise

$K_t = \bar{\Sigma}_t C_t^\top (C_t \bar{\Sigma}_t C_t^\top + Q_t)^{-1}$  Gain

update mean  $\mu_t = \bar{\mu}_t + K_t (z_t - \overset{\text{observation model}}{C_t \bar{\mu}_t})$

update covariance  $\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$  Update



# Simple Multi-dimensional Implementation (also 5 lines of code!)

$$[x \ P] = KF(x, P, z)$$

$$x = A * x;$$


$$P = A * P * A' + Q;$$

$$K = P * C' / (C * P * C' + R);$$

$$x = x + K * (z - C * x);$$

$$P = (\text{eye}(\text{size}(K, 1)) - K * C) * P;$$

# 2D Example



state


$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

measurement

$$\mathbf{z} = \begin{bmatrix} x \\ y \end{bmatrix}$$

Constant position Motion Model

$$\mathbf{x}_t = A\mathbf{x}_{t-1} + B\mathbf{u}_t + \epsilon_t$$



state

$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

measurement

$$\mathbf{z} = \begin{bmatrix} x \\ y \end{bmatrix}$$

Constant position Motion Model

$$\mathbf{x}_t = A\mathbf{x}_{t-1} + B\mathbf{u}_t + \epsilon_t$$

system noise


$$\epsilon_t \sim \mathcal{N}(\mathbf{0}, R)$$

Constant position

$$A = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$B\mathbf{u} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$R = \begin{bmatrix} \sigma_r^2 & 0 \\ 0 & \sigma_r^2 \end{bmatrix}$$



state


$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

measurement

$$\mathbf{z} = \begin{bmatrix} x \\ y \end{bmatrix}$$

Measurement Model

$$\mathbf{z}_t = \mathbf{C}_t \mathbf{x}_t + \delta_t$$



state

$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

measurement

$$\mathbf{z} = \begin{bmatrix} x \\ y \end{bmatrix}$$

Measurement Model

$$\mathbf{z}_t = \mathbf{C}_t \mathbf{x}_t + \delta_t$$

zero-mean measurement noise

$$\mathbf{C} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$\delta_t \sim \mathcal{N}(\mathbf{0}, \mathbf{Q})$$

$$\mathbf{Q} = \begin{bmatrix} \sigma_q^2 & 0 \\ 0 & \sigma_q^2 \end{bmatrix}$$

# Algorithm for the 2D object tracking example



$$A = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

motion model

$$C = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

observation model

## General Case

$$\bar{\mu}_t = A_t \mu_{t-1} + B u_t$$

$$\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^\top + R$$

$$K_t = \bar{\Sigma}_t C_t^\top (C_t \bar{\Sigma}_t C_t^\top + Q_t)^{-1}$$

$$\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$$

$$\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$$

## Constant position Model

$$\bar{\mathbf{x}}_t = \mathbf{x}_{t-1}$$

$$\bar{\Sigma}_t = \Sigma_{t-1} + R$$

$$K_t = \bar{\Sigma}_t (\bar{\Sigma}_t + Q)^{-1}$$

$$\mathbf{x}_t = \bar{\mathbf{x}}_t + K_t (z_t - \bar{\mathbf{x}}_t)$$

$$\Sigma_t = (I - K_t) \bar{\Sigma}_t$$

Just 4 lines of code

```
[x P] = KF_constPos(x, P, z)
```

```
P = P + Q;
```

```
K = P / (P + R);
```

```
x = x + K * (z - x);
```

```
P = (eye(size(K,1)) - K) * P;
```

*Where did the 5th line go?*



## General Case

$$\bar{\mu}_t = A_t \mu_{t-1} + B u_t$$

$$\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^\top + R$$

$$K_t = \bar{\Sigma}_t C_t^\top (C_t \bar{\Sigma}_t C_t^\top + Q_t)^{-1}$$

$$\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$$

$$\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$$

## Constant position Model

$$\bar{\mathbf{x}}_t = \mathbf{x}_{t-1}$$

$$\bar{\Sigma}_t = \Sigma_{t-1} + R$$

$$K_t = \bar{\Sigma}_t (\bar{\Sigma}_t + Q)^{-1}$$

$$\mathbf{x}_t = \bar{\mathbf{x}}_t + K_t (z_t - \bar{\mathbf{x}}_t)$$

$$\Sigma_t = (I - K_t) \bar{\Sigma}_t$$